


Using Instruction tuning to Classify information as privileged



Electronically-Stored Information (ESI) is growing too rapidly to sustain. How can we replace a manual process with an automated/partially-automated (human-in-the-loop) process to handle this data?



OVERVIEW



What is privilege and how can we automate classification



Building the Dataset



Model Selection



Results

PROBLEM



How can we identify privileged information?



What is privilege?

Confidential information/communication that is protected by law (i.e. attorney-client privilege)



Current approach for identifying privileged information requires manually reading/classifying each resource

Not time efficient

Different people may classify the same document differently (subjective)




The amount of data that needs classified grows too quickly to sustain manual classification

DATASET

- Enron Dataset
 - Large, public corpus of approximately 500,000 real-world internal emails from ~150 senior Enron employees, released by the Federal Energy Regulatory Commission (FERC) after the company's 2001 collapse. It is widely used in NLP, social network analysis, and AI training for machine learning.
- Weak Labeling
 - Used to label the original dataset as Privileged/Not Privileged



MODEL DIAGRAM



PREDICTION

- What does the prediction mean?
 - False Positive : Not privileged information classified as 'Privilege'
 - False Negative : Privileged information classified as 'Not Privilege' (legal risk)

		True class (Edgels from the ground truth)	
Predicted class (Edgels from algorithm)	TP (True Positive)	FP (False Positive)	
	FN (False Negative)	TN (True Negative)	

Source: https://www.researchgate.net/figure/A-confusion-matrix-to-define-TP-TN-FP-and-FN_fig3_352393233

EXAMPLE

Input: From: [john.doe@company.com](mailto:john.doe@company.com), To: [jane.smith@company.com](mailto:jane.smith@company.com), Subject: Q3 Budget Review, Body: Hi Jane, please review the attached budget for Q3. Let me know if you have any concerns before Friday.

Prediction: Not Privileged because The budget is for the following year.

Input: From: [legal@company.com](mailto:legal@company.com), To: [ceo@company.com](mailto:ceo@company.com), Subject: Attorney-Client Privileged: Contract Review, Body: This communication is confidential and intended for legal advice. Please review the contract terms regarding liability and indemnification.

Prediction: Privileged because The email is confidential and intended for legal advice.


RESULTS

Experiment #	Hyperparameters Tuned	Accuracy	Precision	Recall	F1 Score
1	None (BERT base)	0.9620	0.8977	0.6311	0.7411
2	Max length 128	0.9470	0.7660	0.5545	0.6433
3	Max length 512	0.9798	0.9558	0.8028	0.8726
4	Train on 15k observations	0.9488	0.8302	0.5104	0.6322
5	Train on 45k observations	0.9714	0.9286	0.7239	0.8136
6	Learning Rate 1e-5	0.9576	0.8897	0.5800	0.7022
7	Learning Rate 3e-5	0.9570	0.8121	0.6520	0.7233
8	Batch Size 16	0.9600	0.8199	0.6868	0.7475
9	Epochs 2	0.9606	0.8634	0.6450	0.7384
10	RoBERTa model	0.9570	0.8803	0.5800	0.6993

WHAT NEXT?




Use a private dataset that includes actual privileged information within a specific field to train



Support a third classification that is "uncertain" that individuals would manually classify



Use RAG for explanation generation with a temperature that discourages creativity (low temperature -> fewer hallucinations)



SUMMARY

- Manually classifying Privileged information is unsustainable
- Privileged data is not publicly available
- Encoder-only model architecture for classification (i.e. BERT)
- Classified emails with an explanation for the prediction

REFERENCES

- 02 - Type of privilege | United States | Global Privilege and Professional Secrecy Guide | Baker McKenzie Resource Hub. (n.d.). <https://resourcehub.bakermckenzie.com/en/resources/global-attorney-client-privilege-guide/north-america/united-states/topics/02---type-of-privilege>
- A BERT based hybrid Recommendation System for Academic Collaboration Vellore Institute of Technology, Chennai, India. (n.d.). <https://arxiv.org/html/2502.15223v1>
- Enron Email Dataset. (n.d.). <https://www.cs.cmu.edu/~enron/>
- Jha, S. (2024). The Future of Contract Review: Innovating Legal Practice through NLP Integration. www.academia.edu. https://www.academia.edu/124591159/The_Future_of_Contract_Review_Innovating_Legal_Practice_Through_NLP_Integration?auto=download
- Maharaj, S. (2025, December 11). The Rise of Categorical Privilege Logs with Growing Popularity and Acceptance. *Consilio*. <https://www.consilio.com/blog/the-rise-of-categorical-privilege-logs-with-growing-popularity-and-acceptance>
- Ratner, A., Bach, S. H., Ehrenberg, H., Fries, J., Wu, S., & Ré, C. (2020). Snorkel: rapid training data creation with weak supervision. *PubMed*. <https://arxiv.org/abs/1711.10160>
- Separating the Privileged Wheat from the Chaff – Using Text Analytics and Machine Learning to Protect Attorney-Client Privilege. (2019, April 24). Richmond Journal of Law and Technology. <https://jolt.richmond.edu/separating-the-privileged-wheat-from-the-chaff-using-text-analytics-and-machine-learning-to-protect-attorney-client-privilege/>
- Statista. (2025, November 19). Data generation volume worldwide 2010-2029. <https://www.statista.com/statistics/871513/worldwide-data-created/>
- Wei, J., Wang, X., Schuurmans, D., Bosma, M., Ichter, B., Xia, F., Chi, E., Le, Q., & Zhou, D. (2022). Chain-of-thought prompting elicits reasoning in large language models. *arXiv*. <https://arxiv.org/abs/2201.11903>
- Zhao, H., Ye, S., & Yang, J. (2021). An Empirical Study on Transfer Learning for Privilege Review. *2021 IEEE International Conference on Big Data (Big Data)*, 37, 2729–2733. <https://doi.org/10.1109/bigdata52589.2021.9672008>